

CSC 4350 – Computer Networks

Fall 2024 Semester

Lecture 5 – The Application Layer/http Protocol

Note

- Material in this lecture is heavily borrowed from Kurose & Ross' "Computer Networking: A Top Down Approach, 8th Edition"

Reliable Data Transfer

- If a protocol provides a guaranteed data delivery service
- One important service that a transport-layer protocol can provide is process-to-process reliable data transfer
- With this service, sending process can just pass its data into the socket and know with confidence that the data will arrive without errors at the receiving process
- Loss-tolerant applications
 - Transport-layer protocol doesn't provide reliable data transfer, some data may not arrive at the receiver
 - Most notably – multimedia applications, such as A/V, conversational A/V that can tolerate a loss
 - Loss of data may result in a small glitch in A/V, but not a “crucial” impairment

Throughput

- Applications that have throughput requirements are called bandwidth-sensitive applications
 - Many current multimedia applications are bandwidth sensitive
 - Some multimedia applications may use adaptive coding techniques to encode digitized voice or video at a rate that matches the currently available throughput
- Elastic applications – specific throughput requirements
 - Can make use of as much or little throughput as is available
 - Examples
 - SMTP/Email
 - FTP
 - Web transfers

Timing

- Transport layer protocol can provide timing guarantees
 - Can be variable
- Example: guarantee that every bit the sender pumps into the socket arrives at the receiver's socket no more than 100 ms later
 - Appealing to interactive real-time applications
 - Internet telephony, virtual environments, conferencing
 - Each provide tight timing constraints on data delivery in order to be effective
- Non-real-time applications – lower delay is always preferable to higher delay, but no tight constraint is placed on end-to-end delays

Security

- Possibly one or more security systems
- Example
 - Sending host – a transport protocol can encrypt all data transmitted by the sending process
 - In the receiving host – transport layer protocol can decrypt the data before delivering the data to the receiving process
 - Provides confidentiality between two processes, even if data is observed in between sender/receiver
- Other ways to achieve confidentiality?

Figure 2.4 – Requirements of Selected Network Applications

Application	Data Loss	Throughput	Time-Sensitive
File transfer/download	No loss	Elastic	No
E-mail	No loss	Elastic	No
Web documents	No loss	Elastic (few kbps)	No
Internet telephony/ Video conferencing	Loss-tolerant	Audio: few kbps–1 Mbps Video: 10 kbps–5 Mbps	Yes: 100s of msec
Streaming stored audio/video	Loss-tolerant	Same as above	Yes: few seconds
Interactive games	Loss-tolerant	Few kbps–10 kbps	Yes: 100s of msec
Smartphone messaging	No loss	Elastic	Yes and no

Transport Services Provided by the Internet

- TCP Services
 - Connection-oriented service
 - Reliable data transfer service
- UDP Services
- Services Not Provided by Internet Transport Protocols

TCP – Connection-Oriented Service

- Has client and server exchange transport-layer control information with each other before the application-level messages start moving
- Alerts the client and server, allowing them to prepare for incoming packets
- After this “handshaking phase,” a TCP connection is said to exist between the sockets of the two processes
 - Full-duplex connection in that the two processes can send messages to each other over the connection at the same time
 - When application finishes sending messages, it must tear down the connection

TCP – Reliable Data Transfer Service

- Communicating processes can rely on TCP to deliver all data sent without error and in the proper order
- When one side of the application passes a stream of bytes into a socket, it can count on TCP to deliver the same stream of bytes to the receiving socket, with no missing/duplicate bytes
- Provides a congestion-control mechanism
 - Throttles a sending process when the network is congested between sender and receiver
 - Attempts to limit each TCP connection to its fair share of network bandwidth

UDP Services

- Lightweight transport protocol, providing minimal services
- Connectionless – no handshaking before the two processes start to communication
- Unreliable data transfer service
 - No guarantee that message sent will ever reach the receiving process
 - Messages may arrive out of order
- No congestion-control mechanism – sending side of UDP can send however much data into the network layer at any rate

Services Not Provided by Internet Transport Protocols

- What about throughput or timing guarantees?
 - Services not provided by today's protocols
- Applications work well-ish because they have been designed to cope with the lack of throughput/timing guarantees
- Figure 2.5 – transport protocols utilized by some popular Internet applications
- TCP gets chosen a bit – reliable data transfer, guaranteeing all data will eventually get to the destination

Figure 2.5 – Popular Internet Applications, Application Layer Protocols, Transport Protocols

Application	Application-Layer Protocol	Underlying Transport Protocol
Electronic mail	SMTP [RFC 5321]	TCP
Remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP 1.1 [RFC 7230]	TCP
File transfer	FTP [RFC 959]	TCP
Streaming multimedia	HTTP (e.g., YouTube), DASH	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary (e.g., Skype)	UDP or TCP

Application-Layer Protocols

- Defines how an application's process, running on different end systems, pass messages to each other
- Particulars
 - Types of messages exchanged; request messages and response messages, for instance
 - Syntax of the various message types, such as fields in the message and how fields are delineated
 - Meaning of the information in the fields
 - Rules for determining when and how a process sends/responds to messages
- Some application-layer protocols are specified in Requests for Comments (RFCs)
 - HTTP – RFC 7230
 - If a browser follows the rules of the HTTP RFC, the browser will be able to retrieve web pages from any web server that has also followed the HTTP RFC
 - Other application-layer protocols are proprietary and not available in the public domain
 - Skype
- A protocol is one piece of a network application

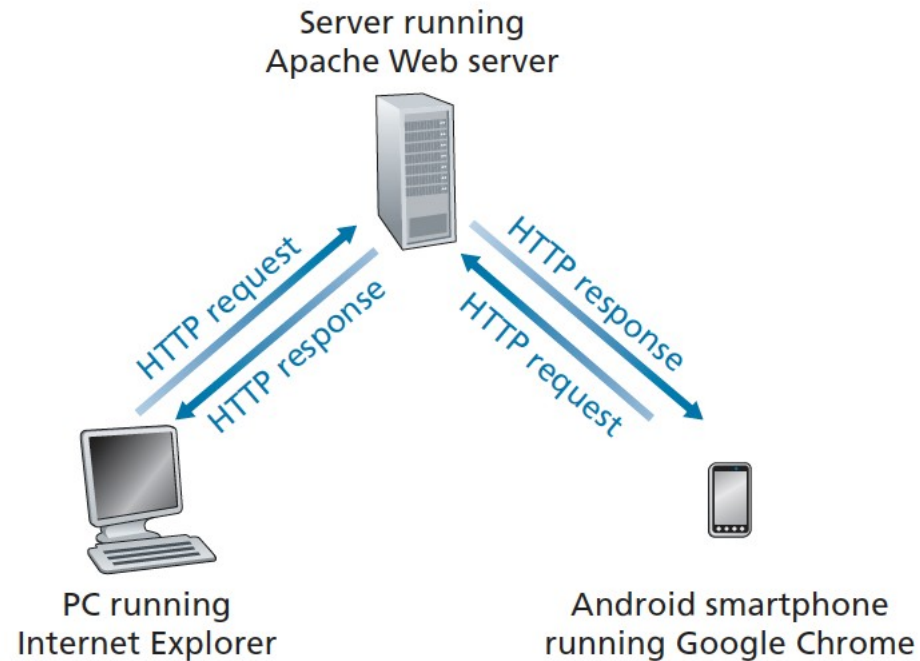
The Web and HTTP

- Overview of HTTP
- Non-Persistent and Persistent Connections
- HTTP Message Formats
- User-Server Interaction – Cookies
- Web Caching
- HTTP/2

Overview of HTTP

- HTTP (HyperText Transfer Protocol) – Web's application-layer protocol – heart of the Internet
- Defined in RFCs 1945, 7230, and 7540
- Implemented with two programs
 - Client
 - Server
 - Interact between each other by exchanging messages
- Web page – consists of objects, usually HTML file, image, CSS, etc.
- Clients – web browsers
- Servers – responsible for handling requests/sending data back to clients

Figure 2.6 – HTTP Request-Response Behavior



Overview of HTTP

- Uses TCP as its transport protocol
 - Client first initiates a TCP connection with the server
 - Once established, the browser and server processes access TCP through socket interfaces
 - Socket
 - Client-side – “door” between client and connection
 - Server-side – “door” between server process and connection
 - Client sends HTTP request messages into its socket interface and receives HTTP messages from its socket interface
 - Server receives requests from socket and sends response messages via socket
 - Once message is sent by client, it is out of client’s hands, and instead with TCP
 - Implication – each HTTP request message sent by a client will eventually reach the server; each server response will eventually be received by the client
- Server sends requested files without storing any state information on the client
 - If client asks for the same object/piece – server will send it again

Non-Persistent and Persistent Connections

- Application developer needs to make a decision wrt to messages going back and forth
 - Should each request/response pair be sent over a separate TCP connection or sent over the same TCP connection
 - Non-persistent connections
 - Persistent connections

HTTP With Non-Persistent Connections

- Steps of transferring a web page from server to client
- Page consists of base HTML file and 10 JPEG images; all 11 objects reside on the same server
- URL
 - <http://www.someSchool.edu/someDepartment/home.index>

Non-Persistent – Steps

- The HTTP client process initiates a TCP connection to the server www.someSchool.edu on port 80 (default port for HTTP)
 - Associated with the TCP connection, there will be a socket at the client and a socket at the server
- The HTTP client sends an HTTP request message to the server via its socket; request message includes the path `/someDepartment/home.index`
- The HTTP server process receives the request message via socket and retrieves the object `/someDepartment/home.index` from its storage
 - Encapsulates the object in an HTTP response message
 - Sends the response message to the client via its socket

Non-Persistent – Steps

- HTTP server process tells TCP to close the TCP connection
 - TCP doesn't actually terminate the connection until it knows for sure that the client has received the message intact
- HTTP client receives the response message
 - TCP connection terminates
 - Message indicates that the encapsulated object is an HTML file
 - Client extracts the file from response message, examines HTML file and finds references to 10 JPEG objects
 - First four steps repeated for each of the referenced JPEG objects

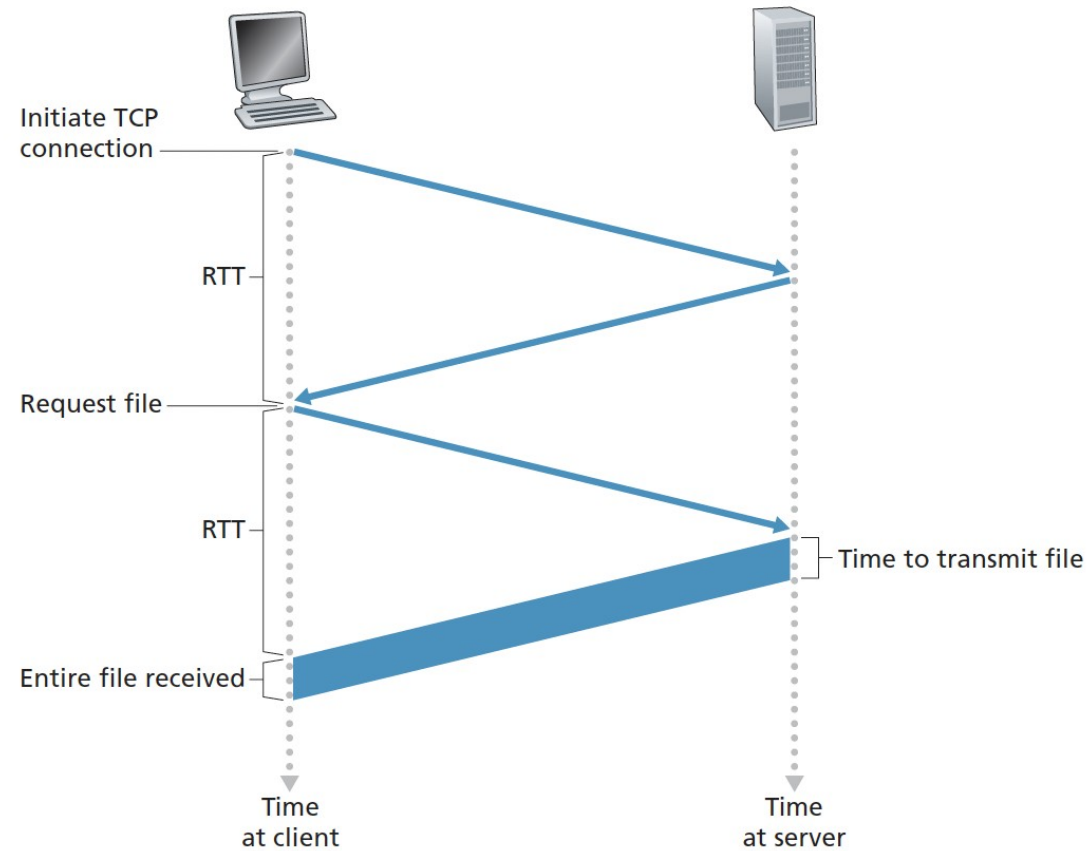
HTTP – Non-Persistent Connections

- Two browsers may interpret a web page in somewhat different ways (not as bad as it used to be)
- HTTP has nothing to do with how a page is interpreted by the client; specifications define only protocol between client and server
- Users can configure browsers to control the degree of parallelism
 - Browsers may open multiple TCP connections and request different parts of the web page over multiple connections

“Back of Envelope” Calculation – Elapsed Time

- Define round-trip time (RTT)
 - Time it takes for a small packet to travel from client to server and then back to the client
 - Includes packet-propagation delays, packet-queuing delays in intermediate routers and switches, and packet-processing delays
 - Shown in Figure 2.7 (next slide)
 - Browser clicks on link, generates TCP connection with server
 - Involves a “three-way” handshake – client sends a small TCP segment to the server
 - Server acknowledges and responds back with a small TCP segment
 - Client acknowledges back to the server
 - First two parts of the three-way handshake take one RTT
 - Then – client sends HTTP request message combined with the third part of the three-way handshake (the acknowledgment)
 - Once request arrives at the server, server sends the HTML file into the TCP connection
 - Requires another RTT
- Response time: $2RTT + \text{transmission time at the server of the HTML file}$

Figure 2.7 – Back-of-the-Envelope Calculation for Time to Request/Receive HTML File



HTTP with Persistent Connections

- Non-persistent connections – issues
 - A brand-new connection must be established and maintained for each requested object
 - For each connection, TCP buffers must be allocated and TCP variables must be kept in both client and server
 - Significant burden – if server handling several requests from different clients all at the same time
 - Also: each object suffers a delivery delay of 2RTT
- HTTP/1.1 persistent connections – server leaves the TCP connection open after sending a response
 - Subsequent requests and responses between the same client and server can be sent over the same connection
 - Multiple web pages residing on the same server can be sent from server to client over a single persistent TCP connection
 - Requests for objects can be made back-to-back without waiting for replies to pending requests (pipelining)

HTTP Message Format – Message

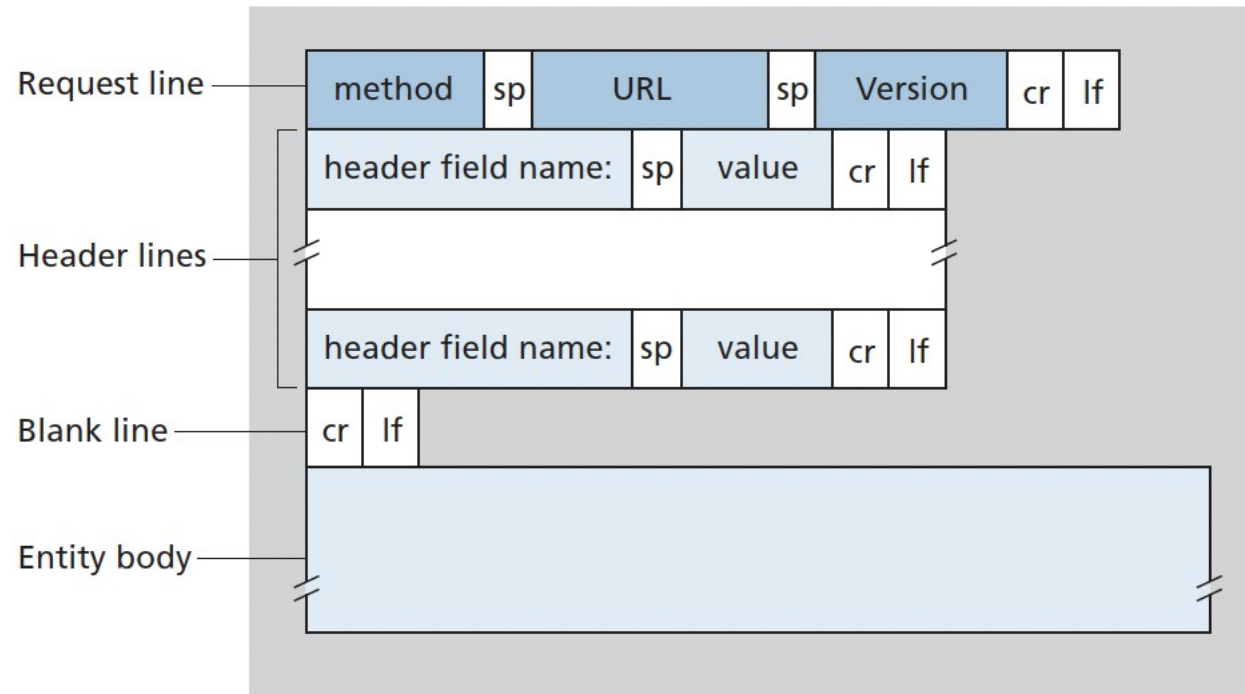
- Message is written in ordinary ASCII – computer-literate human can read it
- Message consists of five lines, each with carriage return and a line feed
- Last line is followed by an additional carriage return and line feed

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
Connection: close
User-agent: Mozilla/5.0
Accept-language: fr
```

HTTP Message Format – Message

- First line of an HTTP request is the request line
 - Three fields
 - Method; possible values: GET, POST, HEAD, PUT, DELETE
 - URL field
 - HTTP version field
- Remaining lines: header lines
 - Host: specifies the host on which the object resides
 - Connection: the browser is telling the server that it doesn't want persistent connections
 - User-agent: browser type making the request (Mozilla/5.0)
 - Accept-language: user prefers to receive a French version of the object, if such an object exists on the server; otherwise it sends its default version
- Majority of HTTP request messages use the GET method
 - Retrieve an object

Figure 2.8 – General Format of an HTTP Request Message



Message Format

- General format closely follows the earlier example
- After header lines, there is an “entire body.”
 - Empty with the GET method
 - Used with the POST method – when the user fills out a form
- The HEAD method is similar to GET
 - When server receives a request, it responds with an HTTP message but leaves out the requested object
 - Developers use HEAD for debugging
- PUT – often used with web publishing tools
 - Allows a user to upload an object to a specific path/directory on a specific web server
 - Also used by applications that need to upload objects to web servers
- DELETE – allows user/application to delete an object on a web server

HTTP Response Message

- Three sections

- An initial status line

- Three fields

- Protocol version
 - Status code
 - Corresponding status message
 - Example: server is using HTTP/1.1 and that everything is okay (object is being sent)

- Six header lines

- Connection: close – tell the client it is going to close the TCP connection after sending the message
 - Date: indicates date and time when HTTP response was created
 - Server – indicates the message was generated by an Apache Web server (similar to User-agent in message request)
 - Last modified – time and date when the object was created/last modified
 - Critical for caching

- Entity body – the meat of the message – the requested object

```
HTTP/1.1 200 OK
Connection: close
Date: Tue, 18 Aug 2015 15:44:04 GMT
Server: Apache/2.2.3 (CentOS)
Last-Modified: Tue, 18 Aug 2015 15:11:03 GMT
Content-Length: 6821
Content-Type: text/html

(data data data data data ...)
```

Figure 2.9 – General Format of an HTTP Response Message

