

# CSC 4350 – Computer Networks

Finish Discussion on Application Layer

Begin Talking about Transport Layer and Its Services

Lecture 10

September 24, 2024

# Note

- Material in this lecture is heavily borrowed from Kurose & Ross' "Computer Networking: A Top Down Approach, 8<sup>th</sup> Edition"

# Socket Programming – Creating Network Applications

- Recall: typical network application consists of a pair of programs – client and server
- Two types of network applications
  - Implementation whose operation is specified in a protocol standard, like an RFC
    - Referred to as “open” since the rules specifying its operation can be known by all
  - Proprietary network application
    - Application-layer protocol that has not been openly published
    - A single developer/developer team creates client and server programs and the developer has complete control over what goes in the code
    - Other independent developers will not be able to write code that works with the application
- Decisions, decisions: TCP or UDP?

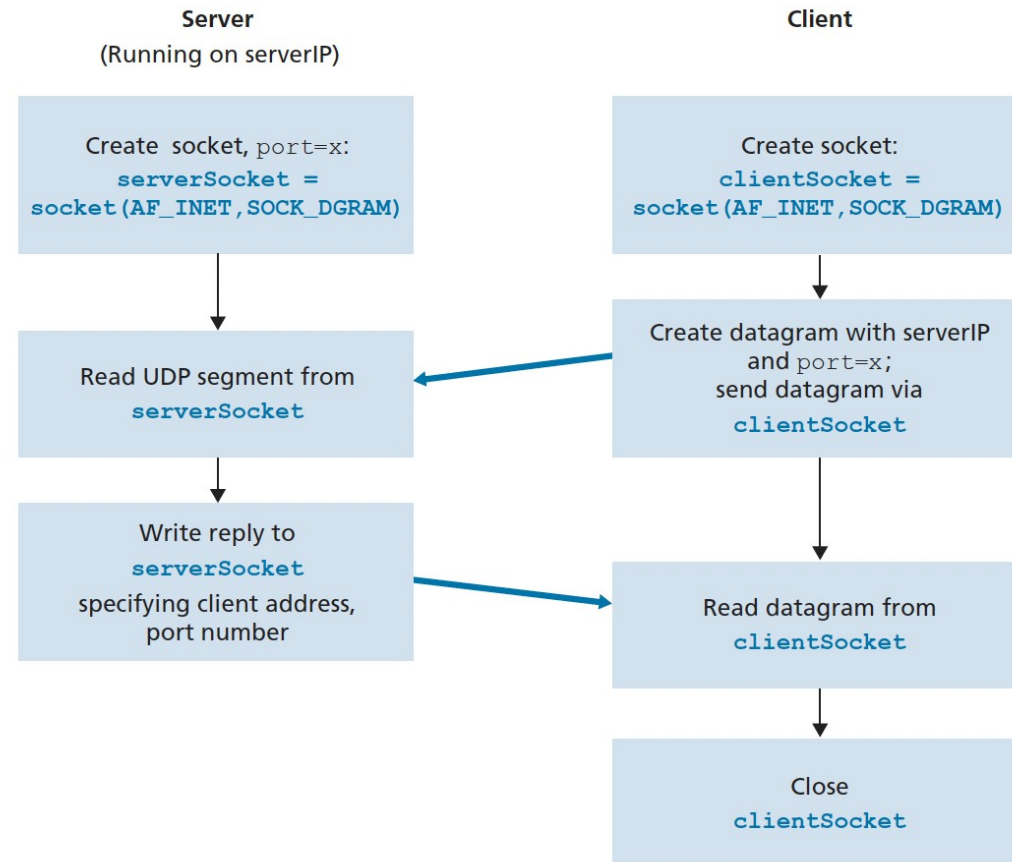
# Closer Look – UDP Sockets

- Before the sending process can push a packet of data out the socket door, when using UDP, it must first attach a destination address to the packet
- After the packet passes through the sender's socket, the internet will use this destination address to route the packet through the internet to the socket in the receiving process
- When the packet arrives at the receiving socket, receiving process will retrieve the packet through the socket and then inspect the packet's contents and take appropriate action

# UDP Sockets

- What goes into the dest address that is attached to a packet?
  - Dest IP address (route packet through internet to destination)
  - Port number is assigned to application, so included in the dest address
  - Sender's IP address of the source host and port number of source socket
  - But – not usually done by the UDP application code; done automatically by the underlying operating system
- Figure 2.27
  - Client reads a line of characters (data) from its keyboard and sends the data to the server
  - Server receives the data and converts the characters to uppercase
  - Server sends the modified data to the client
  - Client receives the modified data and displays the line on its screen

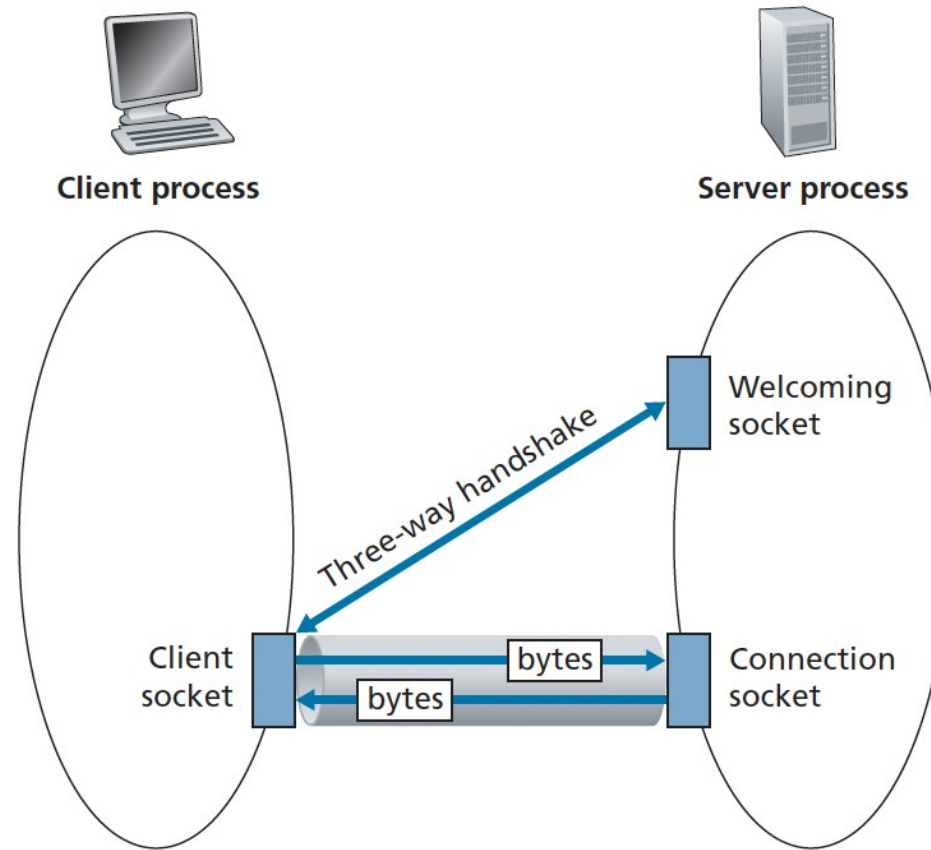
# Figure 2.27 – Client-Server Application Using UDP



# Socket Programming – TCP

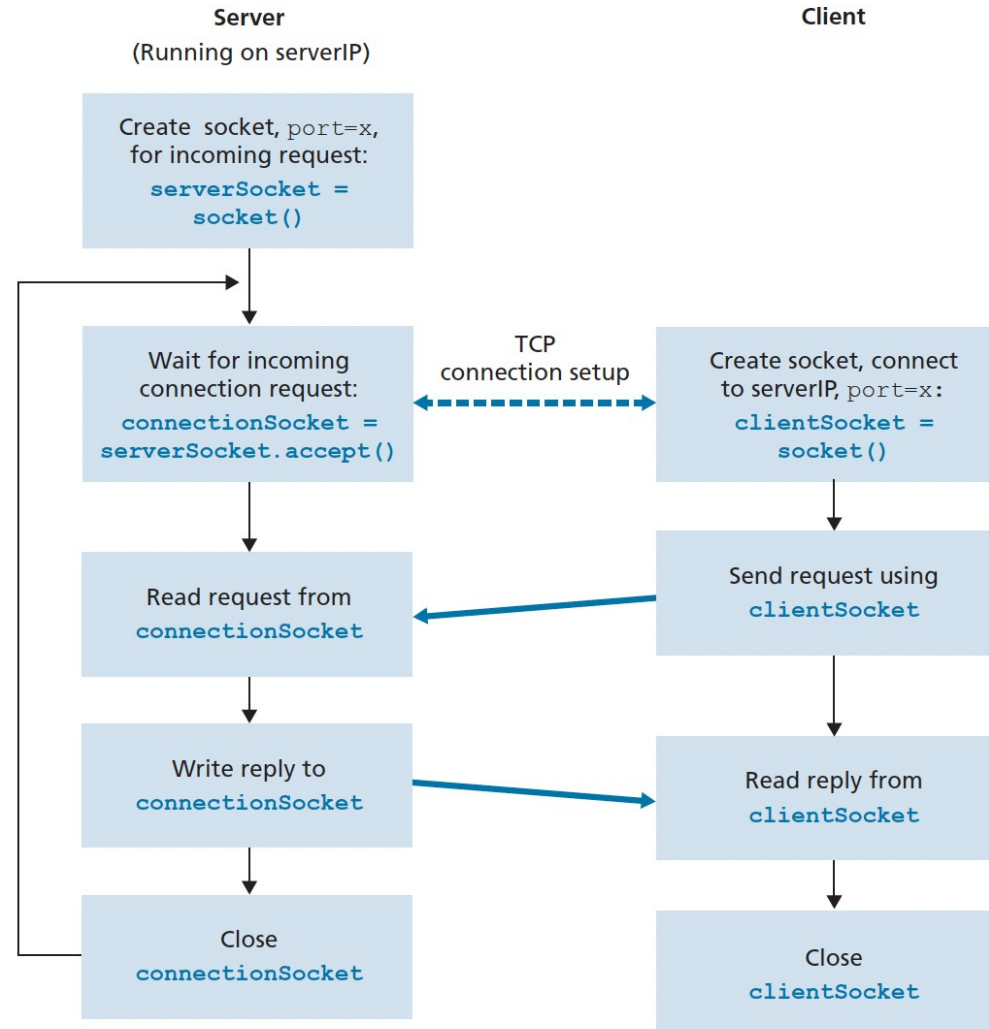
- Before client and server can start to send data to each other, they first need to handshake and establish a TCP connection
  - One end is attached to the client socket
  - Other end – server socket
- When creating the connection, associate it with the client socket address and the server socket address
- When one side wants to send data to the other, data dropped into the TCP connection via socket
  - Slight difference over UDP, where server must attach dest address to packet
- Server has to be ready/must be running as a process before the client attempts to initiate contact
- Server program must have a special socket that welcomes initial contact from a client process running on some host
- Establish connection, three-way handshake, as discussed before for TCP
  - Server will create a new socket dedicated to that particular client
- Application POV – connection socket is connected directly by a pipe

# Figure 2.28 – The TCPServer Process Has 2 Sockets





# Figure 2.29 – Client-Server Application Using TCP



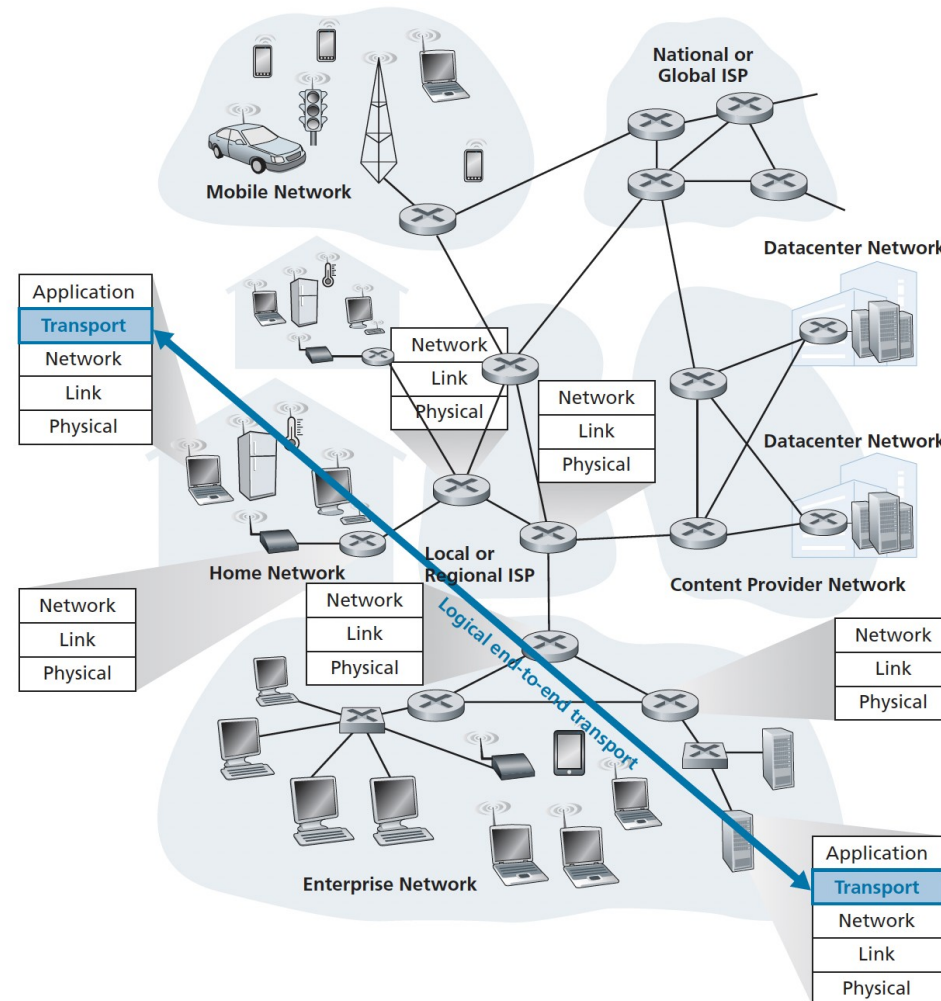
# Transport-Layer Services

- Relationship Between Transport and Network Layers
- Overview of the Transport Layer in the Internet

# Transport-Layer Services - Introduction

- Transport-layer protocol provides for logical communication between application processes running on different hosts
  - “Logical communication” – from an application’s perspective, it is as if the hosts running the processes were directly connected
- Figure 3.1 (next slide)
  - Implemented in the end systems but not in network routers
  - Sender side
    - Converts the application-layer messages it receives from a sending application process to transport-layer packets, known also as segments
    - Done by breaking application messages into smaller chunks and adding a transport-layer header to each chunk to create the segment
    - Passes segment to network layer
    - Network routers act only on the header information; don’t view contents
  - Receiver side
    - Network layer extracts transport-layer segment and passes it up to the transport layer
    - Transport layer then processes the received segment, making the data in the segment available to the receiving application
- Internet has two transport-layer protocols – TCP and UDP

# Figure 3.1 – Transport Layer Providing Logical Rather Than Physical Communication Between Application Processes



# Relationship Between Transport, Network Layers

- Transport-layer protocol provides logical communication between processes running on different hosts, network-layer protocol provides logical-communication between hosts
- East Coast, West Coast houses example in text
  - Application messages = letters in envelopes
  - Processes = cousins
  - Hosts (end systems) = houses
  - Transport-layer protocol = Ann and Bill (distribute mail)
  - Network-layer protocol = postal service (including mail carriers)
- Lead-Up to Talking About TCP and UDP

# Overview of the Transport Layer in the Internet

- UDP – unreliable, connectionless service to the invoking application
- TCP – reliable, connection-oriented service to the invoking application
- When designing a network application, the application developer must specify one of these two protocols
- Separation of terms (for reference elsewhere)
  - TCP – transport layer packet – segment
  - UDP – datagram
- A paragraph later... (your author's notation)
  - Segments – Transport-Layer packet
  - Datagram – Network-Layer packet

# Internet's Network Layer

- Name: IP – Internet Protocol
  - Provides logical communication between hosts
  - Service model: best-effort delivery service
    - IP makes “best effort” to deliver segments between communicating hosts, but no guarantees
      - No guarantee on segment delivery
      - No guarantee on orderly delivery of segments
      - No guarantee on data in the segments
  - Unreliable service
  - Biggest fundamental responsibility for TCP, UDP – extend IP's delivery service between two end systems to a delivery service between two processes running on the end systems
    - Transport layer multiplexing/demultiplexing
  - TCP/UDP provide integrity checking – error detection fields in headers
  - Remember that UDP is not reliable

# Internet's Network Layer

- TCP – several additional services to applications
  - Reliable data transfer – via flow control, sequence numbers, acknowledgments and timers
    - Ensures data gets from sender to receiver
    - Builds off of IP's unreliable service
  - Congestion Control
    - More of a service for the Internet, as a whole
    - Prevents any one TCP connection from swamping the links and routers between hosts with an excessive amount of traffic
    - Tries to give each link an equal share of link bandwidth – regulate the rate at which sending sides of connections can send traffic into the network
    - Not available via UDP



# Multiplexing and Demultiplexing

- Connectionless Multiplexing and Demultiplexing
- Connection-Oriented Multiplexing and Demultiplexing
- Web Servers and TCP

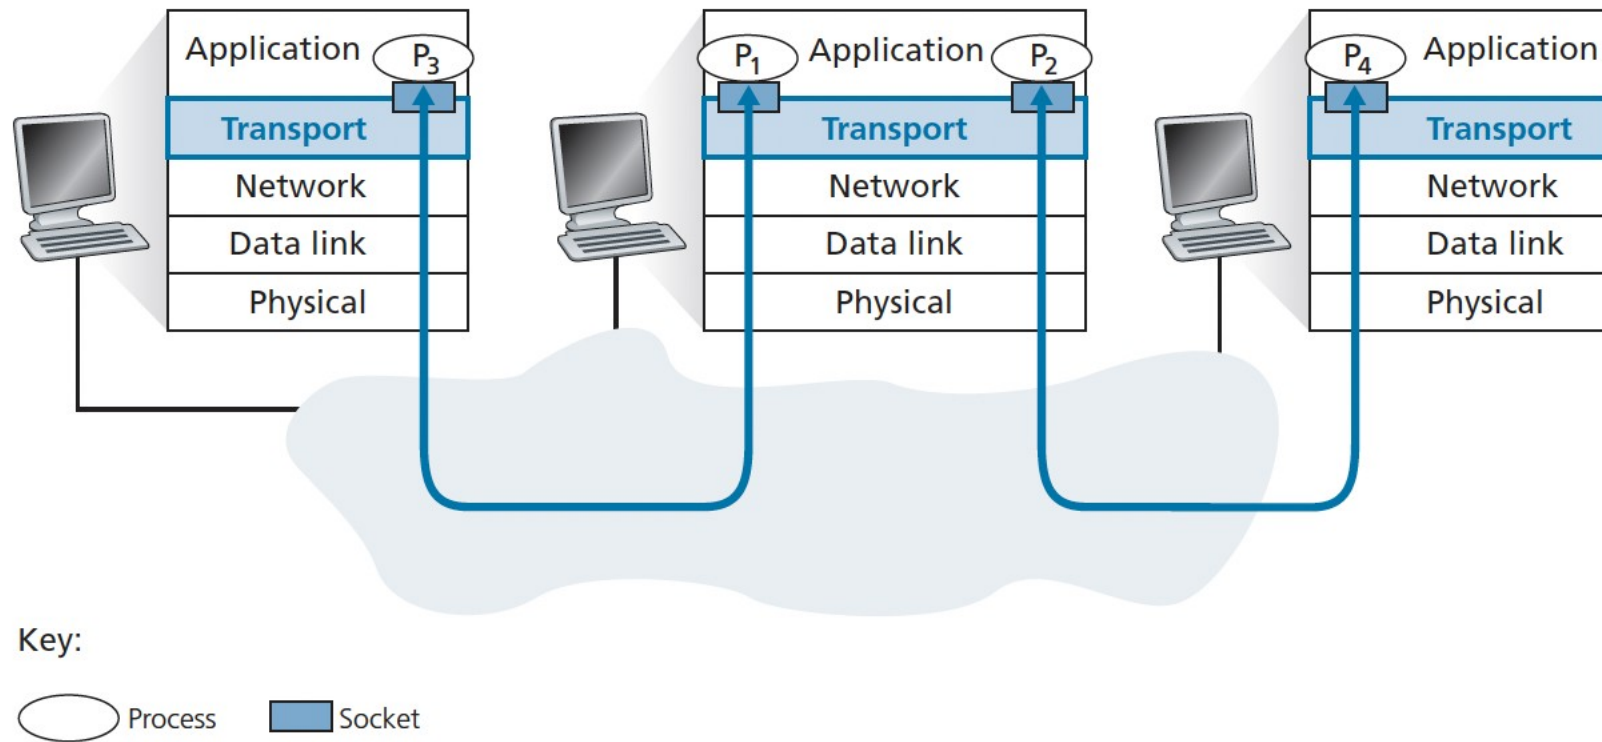
# Multiplexing and Demultiplexing

- Author – keeping discussion on the Internet, but useful for any computer network
- Transport layer – responsibility of delivering the data in segments from the network layer just below
  - Has responsibility of delivering data in these segments to the appropriate application process in the host
  - Example: Using web browser, running ftp, utilizing two Telnet sessions
    - When transport layer in computer receives data from the network layer, it needs to direct the received data to one of these processes
- A process can have 1+ sockets
  - Transport layer in Figure 3.2 does not actually deliver data directly to a process, but to an intermediary socket
  - Because at any given time there can be more than one socket in the receiving host, each socket has a unique identifier
    - Format depends on whether the socket is a UDP or TCP socket

# Multiplexing and Demultiplexing

- Receiving host directs an incoming transport-layer segment to the appropriate socket
  - Each transport-layer segment has a set of fields in the segment for this purpose
  - Receiving end – transport layer examines these fields to identify the receiving socket and then directs the segment to that socket
  - Demultiplexing – job of delivering the data in a transport-layer segment to the correct socket
  - Multiplexing – job of gathering data chunks at the source host from different sockets, encapsulating each data chunk with header information, to create segments, and passing the segments to the network layer

# Figure 3.2 – Transport-Layer Multiplexing/Demultiplexing



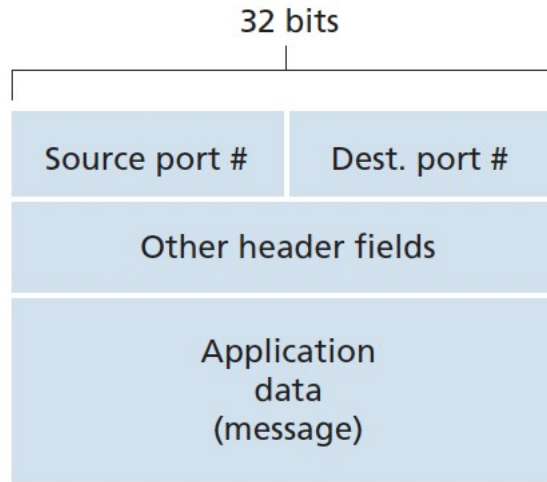
# Multiplexing/Demultiplexing

- Transport layer in middle host – Figure 3.2 – must demultiplex segments arriving from the network layer below either P1 or P2
  - Done by directing the arriving segment's data to the corresponding process' socket
  - Must also gather outgoing data from these sockets, form transport-layer segments, and pass these segments down to the network layer
  - Example: get a bunch of mail at your house
    - Demultiplexing – look at who letters are addressed to, then hand-deliver
    - Multiplexing – collecting letters from others in the house and gives the collected mail to the mail carrier

# Multiplexing/Demultiplexing – Host Perspective

- Figure 3.3
- Transport-layer multiplexing requires
  - Sockets have unique identifiers
  - Each segment has special fields that indicate the socket to which the segment is to be delivered
    - Source port number field
    - Destination port number field
- Each number is a 16-bit number, ranging from 0 to 65535
- Port numbers between 0 to 1023 – well-known port numbers
  - Reserved for use by well-known application protocols, like http, ftp, etc.
- Demultiplexing
  - Each socket in the host could be assigned a port number
  - When a segment arrives at the host, transport layer examines destination port number in the segment, directs segment to the corresponding socket
  - Segment's data passes through the socket into the attached process
  - Above description – how UDP implements...

# Figure 3.3 – Source and Destination Port-Number Fields in a Transport-Layer Segment



# Connectionless Multiplexing and Demultiplexing

- When a UDP socket is created, transport layer automatically assigns a port number to the socket
  - In the range of 1024 to 65535
  - Not currently in use by any other UDP port in the host
- If the application developer writing the code were implementing the server side of a “well-known protocol,” developer would have to assign the corresponding well-known port number
  - Client side – lets transport layer automatically/transparently assign the port number
  - Server side – assigns a specific port number



# UDP Multiplexing/Demultiplexing

- Suppose a process in Host A, with UDP port 19157 wants to send a chunk of application data to a process with UDP port 46428 in Host B
- Transport layer in Host A creates a transport-layer segment that includes the application data, source port number, destination port number, and two other values (not necessary to get into here)
- Transport layer passes the resulting segment to the network layer
- Network layer encapsulates the segment in an IP datagram and makes a best-effort attempt to deliver the segment to the receiving host
- If the segment arrives at the receiving Host B, transport layer at the receiving host examines destination port number and delivers the segment to the appropriate socket
- Note: Host B could be running multiple processes, each with its own UDP socket and associated port number
- As segments arrive, Host B demultiplexes each segment to the appropriate socket by examining the port number

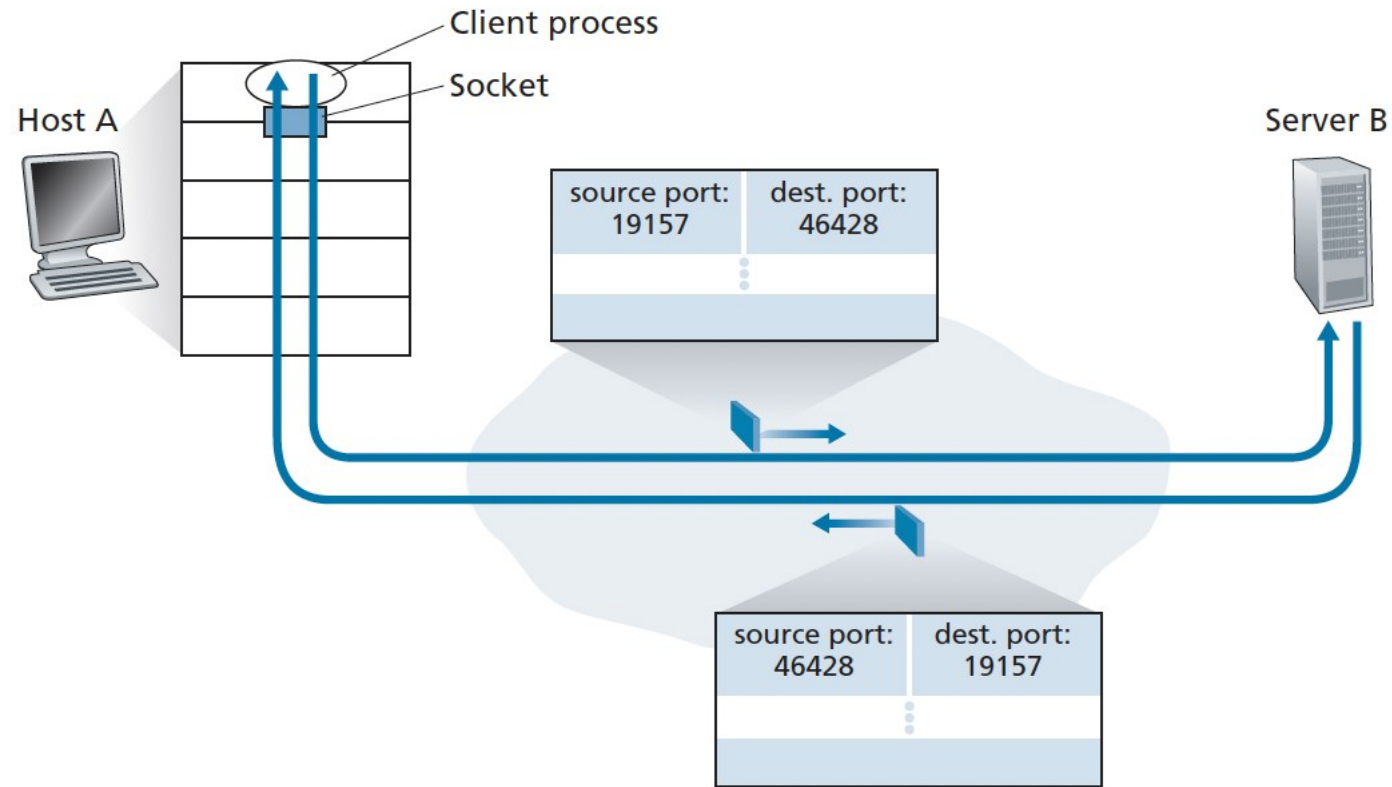
# UDP Multiplexing/Demultiplexing

- UDP socket is fully identified by a two-tuple
  - Destination IP address
  - Destination port number
- If two UDP segments have different IP addresses and/or source port numbers, but have same IP address and destination port number, both will be directed to the same address process via the same socket
- Purpose of source port number – serves as “return address” when B wants to send something back to A

# Connection-Oriented Multiplexing and Demultiplexing

- TCP sockets and TCP connection establishment
  - TCP socket is identified by a four-tuple
    - Source IP address
    - Source port number
    - Destination IP address
    - Destination port number
- When a TCP segment arrives from the network to a host, the host demultiplexes with the four values to send segment to appropriate socket
- Two arriving TCP segments with different source IP addresses or port numbers will be directed to two different sockets
- Server host may support many simultaneous TCP connection sockets
  - Each socket attached to a process and with each socket identified by its own four-tuple
  - When segment arrives at the host, all four fields are used to direct the segment to the appropriate socket

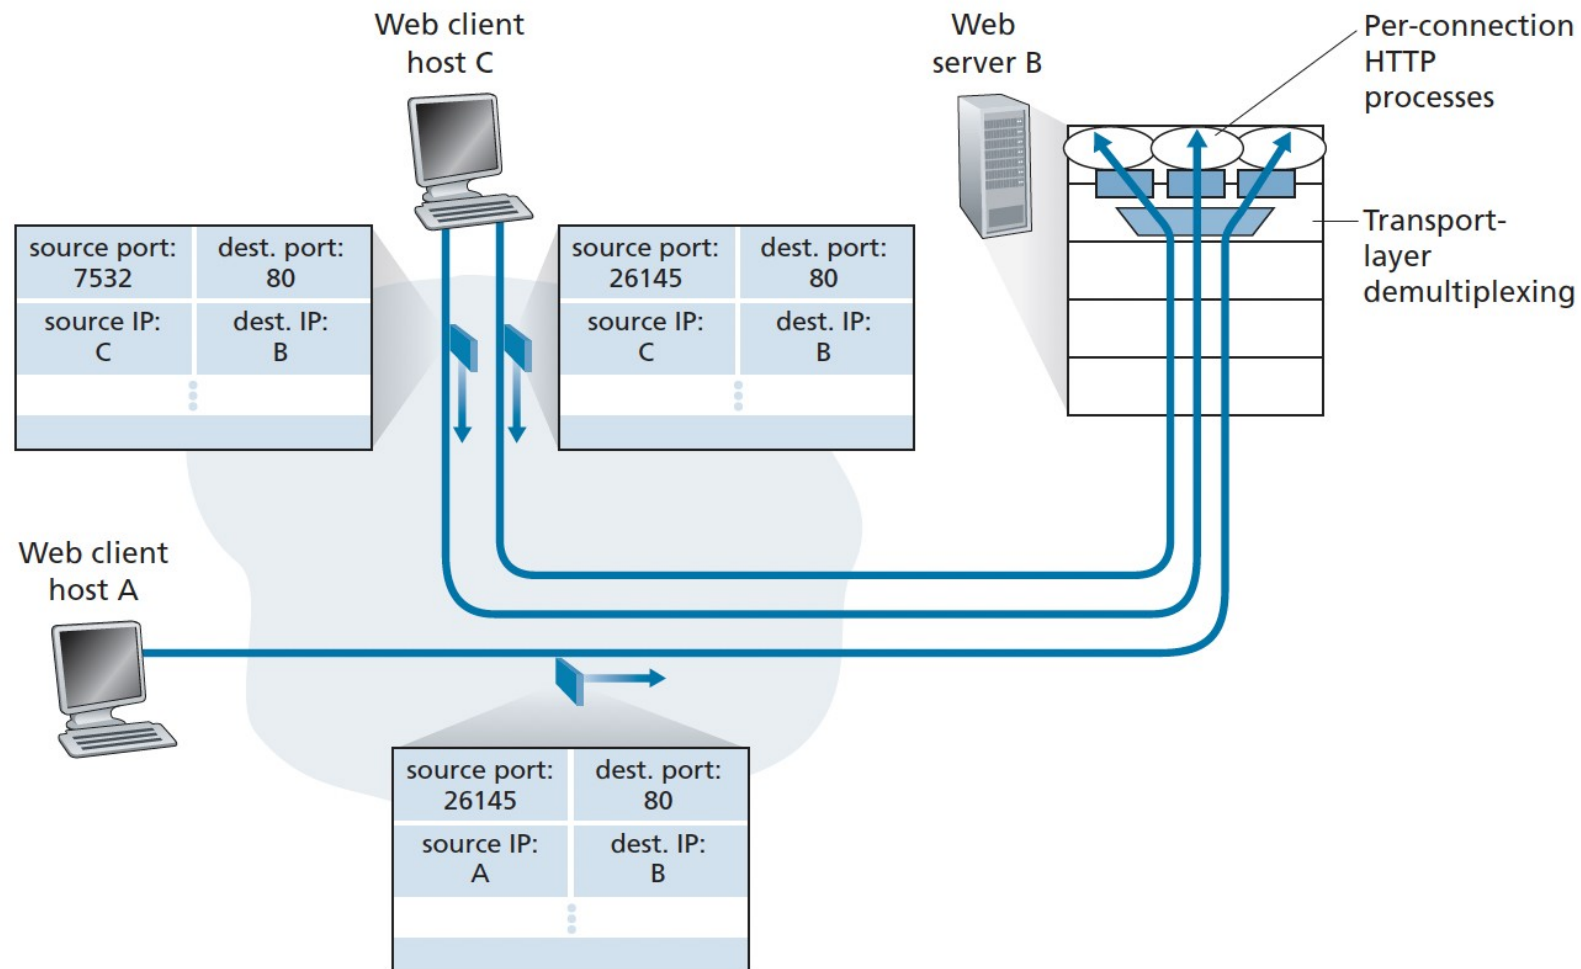
# Figure 3.4 – The Inversion of Source and Destination Port Numbers



# TCP and Figure 3.5

- Host C initiates two http sessions to server B, and Host A initiates one http session to B
- Hosts A and C and server B have their own unique IP address
- Host C assigns two different source port numbers to its two http connections (26145 and 7532)
- Because Host A is choosing source port numbers independently of C, it might also assign a source port of 26145 to its http connection
  - Not an issue – server B will still be able to correctly demultiplex the two connections having the same source port number since the two connections have two different source IP addresses

Figure 3.5 – Two Clients, Same Destination Port Number (80), To Communicate with Same Web Server Application



# Web Servers and TCP

- Host runs on a web server, like an Apache web server, on port 80
  - When clients send segments to the server, all segments will have destination port 80
  - Both the initial connection-establishment segments and segments carrying http request messages will have destination port 80
  - Different clients will have different source IP addresses and port numbers
- Figure 3.5
  - Web server that creates a new process for each connection
  - Each of these processes has its own connection socket through which http requests arrive and responses sent
  - There is not always a 1:1 correspondence between connection sockets and processes
    - Web servers often use only one process and create a new thread with a new connection socket for each client connection

# Figure 3.5 Discussion – Continued

- If client and server are using persistent http, then throughout the duration of the persistent connection the client and server exchange http messages via same server socket
- If non-persistent http
  - A new socket is created and later closed for every request/response and a new socket is created and closed for every request/response
  - Can impact performance of a busy web server